

IBM Resilient



Incident Response Platform Integrations

Cloud Foundry Function V1.0.0

Release Date: August 2018

Resilient Functions simplify development of integrations by wrapping each activity into an individual workflow component. These components can be easily installed, then used and combined in Resilient [Platform](#) workflows. The Resilient platform sends data to the function component that performs an activity then returns the results to the workflow. The results can be acted upon by scripts, rules, and workflow decision points to dynamically orchestrate the security incident response activities.

This guide describes the Cloud Foundry Function.

Overview

This function interfaces with Cloud Foundry platform to allow user to manage deployed applications, their instances, and deploy new applications. Managing the applications includes starting/stopping, updating, restaging, deleting, and getting various types of information about it. All of these possibilities are separated into 3 functions that can be used in the workflows.

It is a wrapper around Cloud Foundry's API with possibilities to be adjusted for different platform providers. On the moment of release the latest stable version is: <https://apidocs.cloudfoundry.org/3.1.0/>

Installation

Before installing, verify that your environment meets the following prerequisites:

- Resilient platform is version 30 or later.
- You have the necessary credentials to access your provider's Cloud Foundry API
- You have a Resilient account to use for the integrations. This can be any account that has the permission to view and modify administrator and customization settings, and read and update incidents. You need to know the account username and password.
- You have access to the command line of the Resilient appliance, which hosts the Resilient platform; or to a separate integration server where you will deploy and run the functions code. If using a separate integration server, you must install Python version 2.7.10 or later, or version 3.6 or later, and "pip". (The Resilient appliance is preconfigured with a suitable version of Python.)

Formatted: Font: Not Italic, Font color: Auto

Deleted: <Overview of solution which this integration utilizes>¶
<Description of the functional operations (i.e. These functions create ..., search ..., collect ..., enforce ..., etc.>¶
<EXAMPLE: The Utility Functions integration package contains several useful workflow functions for common automation and integration activities in the Resilient platform.>¶
<The remainder of this document describes each included function, how to configure them in custom workflows, and any additional customization options.>¶

Licensed Materials – Property of IBM

© Copyright IBM Corp. 2010, 2018. All Rights Reserved.

US Government Users Restricted Rights – Use, duplication or disclosure restricted by GSA ADP Schedule Contract with IBM Corp.

Install the Python components

The functions package contains Python components that are called by the Resilient platform to execute the functions during your workflows. These components run in the Resilient Circuits integration framework.

The package also includes Resilient customizations that will be imported into the platform later.

Complete the following steps to install the Python components:

1. Ensure that the environment is up-to-date, as follows:

```
sudo pip install --upgrade pip
sudo pip install --upgrade setuptools
sudo pip install --upgrade resilient-circuits
```

2. Run the following command to install the package:

```
sudo pip install --upgrade fn_cloud_foundry-1.0.0.tar.gz
```

Configure the Python components

The Resilient Circuits components run as an unprivileged user, typically named integration. If you do not already have an integration user configured on your appliance, create it now.

Complete the following steps to configure and run the integration:

1. Using sudo, switch to the integration user, as follows:

```
sudo su - integration
```

2. Use one of the following commands to create or update the resilient-circuits configuration file. Use -c for new environments or -u for existing environments.

```
resilient-circuits config -c
```

or

```
resilient-circuits config -u
```

3. Edit the resilient-circuits configuration file, as follows:

- a. In the [resilient] section, ensure that you provide all the information required to connect to the Resilient platform.

- b. In the [fn cloud foundry] section, edit the settings as follows:

```
[fn cloud foundry]
#Base url endpoint of your CF platform
#For example, for IBM's BlueMix it is: https://api.stag1.ng.bluemix.net/
cf api base=xxx
#Enter only what's required by your authenticator.
#For example, the default BlueMixCF authenticator only requires apikey.
cf api apikey=xxx
cf api username=xxx
cf api password=xxx
```

4. Current function is configured by default to work with BlueMix Cloud Foundry platform, but it can be connected to any platform that exposes API.

Deleted: fn_<fn_name>

Deleted: <version>.<zip>

Deleted: <If this is a zip package with a tar.gz file inside, change step 2 to this:¶
To install the package, you must first unzip it then install the package as follows:¶
sudo pip install --upgrade fn_<fn_name>-<version>.<tar.gz>¶

Deleted: fn_<fn_name>

Deleted: <settings>

Formatted: Body Text, Numbered + Level: 1 + Numbering Style: 1, 2, 3, ... + Start at: 1 + Alignment: Left + Aligned at: 0" + Indent at: 0.25"

- a. Create a new **Authenticator** class, that inherits from **AuthenticatorBase** located in `utils/authentication`. It needs to implement 2 methods, `auth` – to authenticate, and `get_headers` – to return headers that need to be added to the http requests in order for it to be authenticated as a dictionary.
- b. In the code for the functions, the default authenticator will need to be replaced by the one you've created.
- c. Build the package with the new code, by running `python setup.py sdist`, and install the newly built package.

Formatted: Body Text, Numbered + Level: 2 + Numbering Style: a, b, c, ... + Start at: 1 + Alignment: Left + Aligned at: 0.5" + Indent at: 0.75"

Deploy customizations to the Resilient platform

The package contains 3 functions that can be used in workflows, as well as an example workflow that demonstrates how to call one of the functions.

1. Use the following command to deploy these customizations to the Resilient platform:

```
resilient-circuits customize
```

2. Respond to the prompts to deploy functions, message destinations, workflows and rules.

Formatted: Font: Not Italic, Font color: Auto

Deleted: <Describe what the package contains.>¶
<EXAMPLE: The package contains function definitions that you can use in workflows, and includes example workflows and rules tha show how to use these functions.>¶

Run the integration framework

To test the integration package before running it in a production environment, you must run the integration manually with the following command:

```
resilient-circuits run
```

The resilient-circuits command starts, loads its components, and continues to run until interrupted. If it stops immediately with an error message, check your configuration values and retry.

Configure Resilient Circuits for restart

For normal operation, Resilient Circuits must run continuously. The recommend way to do this is to configure it to automatically run at startup. On a Red Hat appliance, this is done using a systemd unit file such as the one below. You may need to change the paths to your working directory and `app.config`.

1. The unit file must be named `resilient_circuits.service` To create the file, enter the following command:

```
sudo vi /etc/systemd/system/resilient_circuits.service
```

2. Add the following contents to the file and change as necessary:

```
[Unit]
Description=Resilient-Circuits Service
After=resilient.service
Requires=resilient.service

[Service]
Type=simple
User=integration
WorkingDirectory=/home/integration
ExecStart=/usr/local/bin/resilient-circuits run
Restart=always
TimeoutSec=10
Environment=APP_CONFIG_FILE=/home/integration/.resilient/app.config
```

Deleted: <replace the contents below with your own>

```
Environment=APP_LOCK_FILE=/home/integration/.resilient/resilient_circuits.  
lock  
[Install]  
WantedBy=multi-user.target
```

3. Ensure that the service unit file is correctly permissioned, as follows:

```
sudo chmod 664 /etc/systemd/system/resilient_circuits.service
```

4. Use the systemctl command to manually start, stop, restart and return status on the service:

```
sudo systemctl resilient_circuits [start|stop|restart|status]
```

You can view log files for systemd and the resilient-circuits service using the journalctl command, as follows:

```
sudo journalctl -u resilient_circuits --since "2 hours ago"
```

Function Descriptions

Once the function package deploys the function(s), you can view them in the Resilient platform Functions tab, as shown below. This package also includes an example workflow showing how to use one of the functions. You can edit and modify this workflow to your needs.

Name	Description
Cloud Foundry: Application Command	Performs a specified action on the chosen Cloud Foundry applications. Currently available actions: start, stop, restage, info, delete, inst
Cloud Foundry: Create an App	Creates and deploys a cloud foundry applications from the specified parameters/docker files. additional_parameters accept any addition that CF API accepts in JSON format. Format of the results is described in the doc.
Cloud Foundry: Instance Command	Performs an action on the specified instances of an application. Currently available: delete. Multiple instances can be provided, but only name should be given. Format of the result is written in the doc.

Deleted: <IF NEEDED - The package also includes example workflows and rules that show how the functions can be used. You can copy and modify these workflows and rules for your own needs.>

fn_cloud_foundry_manage_applications; Cloud Foundry: Application Command

Name *	Cloud Foundry: Application Command
API Name *	fn_cloud_foundry_manage_applications
Message Destination *	cloud_foundry
Description	Performs a specified action on the chosen Cloud Foundry applicati Currently available actions: start, stop, restage, info, delete, instanc Format of the response is described in the doc.

Inputs

fn_cloud_foundry_action

fn_cloud_foundry_applications

fn_cloud_foundry_additional_parameters_json

Deleted: Screenshot – Function tab or individual function¶
<NOTES:¶
If you have only one function, you can use a screenshot of the function itself, which you get by clicking the function name in the Functions tab. Then change the text above to reference the right screenshot.¶
If you have multiple functions, use a Heading 2 for each function as shown below>¶
Fn_<name>

This function performs one of the chosen actions on the specified Cloud Foundry applications.

It takes in 3 inputs:

- fn_cloud_foundry_applications: **Required** – name or comma-separated names of the application(s) that the actions should be applied to.
- fn_cloud_foundry_action: **Required** – select field with the list of all possible actions to be performed on the application. They include start, stop, restage, delete, instances, and info.
- fn_cloud_foundry_additional_parameters_json: **Optional** – here for the purposes of allowing the actions to be added and extended in the future. Accepts valid JSON, which will be passed to the method executing chosen actions.

Formatted: Bulleted + Level: 1 + Aligned at: 0.25" + Indent at: 0.5"

Listed below, is the format of the results that the function produces. It would list whether the action was successfully applied to each of the applications specified, as well as other data relevant to the action itself, per application.

```
{
  "application-1": {
    "success": true,
    "other_data": "...",
    "keys": ["success", "other_data"]
  },
  "application-2": {
    "success": false,
    "details": "...",
    "keys": ["success", "details"]
  },
  "keys": ["application-1", "application-2"]
}
OR, if the action is incorrect:
{
  "success": false,
  "details": "...",
  "keys": ["success", "details"]
}
```

"keys" store the list of keys for each level of a dictionary to provide convenience in post-processing.

Deleted: <function display name in Resilient Functions tab>

Deleted: <Describe the function.>
<Describe any actions the user needs to take on the Resilient platform, such as adding inputs or creating a field for Layouts.>
Describe any other considerations for the supplied rules and configurations needed for workflows. Think about parameters which need to be modified per each customer's environment.
<If there are scenarios of use which would help a customer understand how to use the function, indicate them here.>
<EXAMPLE: This function produces hashes of a file attached to an incident. Provide the incident ID and attachment ID as input. The output includes md5, sha1, sha256 and other hashes of the file content. Those hashes can then be used as artifacts or in other parts of your workflows.>

Formatted: Code

Formatted: Don't keep with next

fn_cloud_foundry_create_app: Cloud Foundry: Create an App

Name *	Cloud Foundry: Create an App
API Name * ⓘ	fn_cloud_foundry_create_app
Message Destination *	cloud_foundry
Description	Creates and deploys a cloud foundry applications from the specific parameters/docker files. additional_parameters accept any additional parameters that CF API accepts in JSON format. Format of the res

Inputs

fn_cloud_foundry_applications

fn_cloud_foundry_space_guid

fn_cloud_foundry_additional_parameters_json

This function creates a new application in Cloud Foundry's platform.

It takes in 3 inputs:

- fn_cloud_foundry_applications: **Required** – single name of the app to be created.
- fn_cloud_foundry_space_guid: **Required** – GUID of the space in which the app should be created.
- fn_cloud_foundry_additional_parameters_json: **Optional** –Accepts valid JSON, which will be passed to the method executing chosen actions.

Can be used to specify the fields needed for the creation of the **Docker** application, as stated in the API.

Return value will contain the information about the app, as provided in the API as well as "success" field to determine whether the app was, in fact, created.

Formatted: Indent: Left: 0.5", No bullets or numbering

Formatted: No bullets or numbering

fn_cloud_foundry_instance_command: Cloud Foundry: Instance Command

Formatted: Heading2, No bullets or numbering, Don't keep with next

Name *	Cloud Foundry: Instance Command
API Name * ⓘ	fn_cloud_foundry_instance_command
Message Destination *	cloud_foundry
Description	Performs an action on the specified instances of an application. Currently only available: delete. Multiple instances can be provided, but only one application name should be given. Format of the result is written in

Inputs

fn_cloud_foundry_instance_action

fn_cloud_foundry_instances

fn_cloud_foundry_applications

fn_cloud_foundry_additional_parameters_json

Function that perform commands on the instances of a specified application. Currently only supports deleting an instance.

It takes in 4 inputs:

- fn_cloud_foundry_instances: **Required** – name or comma-separated names of the instances(s) that the actions should be applied to.
- fn_cloud_foundry_applications: **Required** – name of the single application whose instances should be affected by the action.
- fn_cloud_foundry_instance_action: **Required** – select field with the list of all possible actions to be performed on the instance. Currently, only delete.
- fn_cloud_foundry_additional_parameters_json: **Optional** – here for the purposes of allowing the actions to be added and extended in the future. Accepts valid JSON, which will be passed to the method executing chosen actions.

Troubleshooting

There are several ways to verify the successful operation of a function.

- Resilient Action Status

When viewing an incident, use the Actions menu to view Action Status. By default, pending and errors are displayed. Modify the filter for actions to also show Completed actions. Clicking on an action displays additional information on the progress made or what error occurred.

- Resilient Scripting Log

A separate log file is available to review scripting errors. This is useful when issues occur in the pre-processing or post-processing scripts. The default location for this log file is: `/var/log/resilient-scripting/resilient-scripting.log`.

- Resilient Logs

By default, Resilient logs are retained at `/usr/share/co3/logs`. The `client.log` may contain additional information regarding the execution of functions.

- Resilient-Circuits

The log is controlled in the `.resilient/app.config` file under the section `[resilient]` and the property `logdir`. The default file name is `app.log`. Each function will create progress information. Failures will show up as errors and may contain python trace statements.

Support

For additional support, contact support@resilientsystems.com.

Including relevant information from the log files will help us resolve your issue.

Formatted: No bullets or numbering, Don't keep with next

Deleted: ¶

Resilient Platform Configuration¶

<Only use this section if you have general configuration advice or requirements, such as describing issues moving the configuration from a testing environment to a production environment. Think about the parameters that need to be modified per each customer's environment.>¶

<If your function exports data tables or custom fields, the changes in the Resilient platform layouts may not be preserved. Therefore, you should document those changes in this section and advise the reader to manually make those changes in the platform.>¶

<If there are scenarios of use which would help a customer understand how to use the Integration as a whole (not per function), indicate them here.>¶

<If not needed, delete this section.>¶

Formatted: Font: Not Italic, Font color: Auto

Deleted: ¶

Documentation Guidelines¶

<Do NOT include this section in your guide.>¶

Here are some writing guidelines:¶

Never use "Resilient," instead use "Resilient platform."¶

Use "deploy to the Resilient platform" to describe the resilient-circuits customize command.¶

Do not initial cap function, workflow, etc. unless you are referring to a specific item (proper name), such as

Utilities Function.¶

Try to avoid passive voice and future tense.¶

For the guide's file name, use this format: Resilient

Integration <name> Function¶

In the Word file, open properties and make these

changes:¶

Author = IBM Resilient¶

Title = Resilient IRP Integrations <name> Function

Guide¶

If you don't know how to open properties:¶

Click **File > Info**.¶

On the right side of the Info page, click the **Properties**

drop-down and select **Show Document Panel**. This

shows the Document Properties with the Author and

Title fields.¶